



Boundary conditions on non-planar boundaries

Louis Moresi¹ 

¹Australian National University

Keywords Underworld Code, Tricks of the Trade, development

A free surface moves under the traction it carries. That makes the wall-normal stress the quantity driving the model rather than something read off at the end, and it is the reason we went back to how the boundary condition underneath it is imposed.

The condition itself is ordinary. Free slip is no flow through the boundary and no tangential drag along it,

$$\mathbf{u} \cdot \hat{\mathbf{n}} = 0 \quad \text{and} \quad \hat{\mathbf{t}} \cdot \boldsymbol{\sigma} \hat{\mathbf{n}} = 0. \quad (1)$$

On a box the first of those is a single velocity component. You hold u_x on a vertical wall, the solver removes a row, and there is nothing further to discuss. On a sphere, an annulus, a mesh that has been moved, or a surface with topography, $\mathbf{u} \cdot \hat{\mathbf{n}}$ is not a component of anything. That is the whole difficulty, and everything below is a way around it.

The second condition is worth naming because it is the one people forget. Zero tangential traction is *natural*: it is what you get by leaving the boundary term out of the weak form. Nothing has to be done to impose it, and something has to be done to avoid imposing it accidentally.

1. WHERE THE BOUNDARY TERM COMES FROM

Every method below is a statement about one term. Multiplying the momentum balance by a test function $\mathbf{w}$ and integrating by parts gives

$$\int_{\Omega} \boldsymbol{\sigma} : \nabla \mathbf{w} \, dV - \int_{\partial\Omega} (\boldsymbol{\sigma} \hat{\mathbf{n}}) \cdot \mathbf{w} \, dS = \int_{\Omega} \mathbf{f} \cdot \mathbf{w} \, dV. \quad (2)$$

Drop the surface integral and you have imposed zero traction in both directions — free *everything*, not free slip. Free slip keeps the tangential half of that and replaces the normal half with the constraint. How you do the replacing is the choice.

2. FOUR WAYS, IN ORDER

2.a. A direct penalty

Add a term that punishes any flow through the boundary:

$$\dots + \frac{\gamma}{h} \int_{\partial\Omega} (\mathbf{u} \cdot \hat{\mathbf{n}})(\mathbf{w} \cdot \hat{\mathbf{n}}) \, dS. \quad (3)$$

One line, no new machinery, and it works on any geometry. It is still in a good many working scripts, and deservedly. What you are solving is a perturbed problem, though, and it is perturbed by exactly the amount the constraint is violated: the discrete solution sits where the penalty term balances the traction it is fighting, which leaves $\mathbf{u} \cdot \hat{\mathbf{n}}$ small but not zero. Making it smaller means pushing harder, and pushing harder conditions the operator worse. The error is traded against the conditioning, and measured below, that trade runs out: past 10^4 the leak stops improving, and by 10^6 the solve fails. Nitsche moves the floor down rather than removing it — it fails too, at its own threshold.

Underworld spells it as a boundary traction opposing normal flow, using the facet normal at the quadrature points:

```
G = mesh.Gamma
stokes.add_natural_bc(1.0e4 * G.dot(v.sym) * G, "Upper")
```

2.b. Nitsche

The reason the penalty is only accurate in the limit is that it is not *consistent*: substituting the true solution does not leave the equation satisfied, because the true solution is subject to a boundary traction the penalty form ignores. Nitsche's method (Nitsche, 1971) restores consistency by carrying that traction explicitly:

$$\dots - \int_{\partial\Omega} (\hat{\mathbf{n}} \cdot \boldsymbol{\sigma}(\mathbf{u})\hat{\mathbf{n}})(\mathbf{w} \cdot \hat{\mathbf{n}}) \, dS - \int_{\partial\Omega} (\hat{\mathbf{n}} \cdot \boldsymbol{\sigma}(\mathbf{w})\hat{\mathbf{n}})(\mathbf{u} \cdot \hat{\mathbf{n}}) \, dS + \frac{\gamma}{h} \int_{\partial\Omega} (\mathbf{u} \cdot \hat{\mathbf{n}})(\mathbf{w} \cdot \hat{\mathbf{n}}) \, dS. \quad (4)$$

The first of the three is the consistency term: it is the boundary traction the integration by parts produced, and putting it back is what makes the true solution satisfy the discrete equations exactly. The second is its transpose, which keeps the form symmetric and buys optimal convergence in L^2 . The third is the penalty again, and it is still needed — but now for *stability* rather than for accuracy, and γ has a threshold set by an inverse inequality rather than being a dial you turn up until the answer looks right.

This is a real improvement and it is still a weak imposition. The constraint holds to the accuracy of the discretisation, not to the accuracy of the arithmetic — measured below, it leaks a few parts in a thousand on a mesh of the resolution people actually run, and the leak falls with the mesh rather than with the machine.

2.c. A constraint equation, with a multiplier

The two above add a *term* to an equation. This one adds an *equation*.

Carry a scalar field h on the boundary and require, as a row of the system in its own right,

$$\int_{\partial\Omega} (\mathbf{u} \cdot \hat{\mathbf{n}} - g) q \, dS = 0 \quad \text{for all } q, \quad (5)$$

with h entering the momentum row as the traction $h\hat{\mathbf{n}}$ that holds the constraint. It is a Lagrange multiplier, and the system becomes a larger saddle point: velocity, pressure, and now h .

The constraint row is exact, so unlike a penalty there is no parameter whose size decides how well it holds. Two practical things do have to be dealt with.

- h is carried as a full-domain field but only its boundary trace means anything, so the interior degrees of freedom are constrained out of the global system in the section before the solve sees it. They are not solved for and they do not enlarge the $[p, h]$ block. The interior rows are the screening block alone, so the interior multiplier is determined by the boundary trace and carries no independent signal.
- The $[p, h]$ Schur complement is poorly conditioned on its own, so an augmented-Lagrangian term $r(\mathbf{u} \cdot \hat{\mathbf{n}} - g)\hat{\mathbf{n}}$ is added to the momentum row. It does not change what the constraint enforces, because the h row still carries the exact constraint. It does change what h is, which is the subject of a section below.

```
stokes = uw.systems.Stokes_Constrained(mesh, velocityField=v, pressureField=p)
h = stokes.add_constraint_bc(0.0, "Upper")
stokes.solve()
```

And the reason to care about it beyond the constraint: **at convergence the momentum row's boundary term is the normal traction**. It is not recovered from the velocity field afterwards — it is an unknown the solve returned, available through `multiplier` and, divided by $\Delta\rho g$, through `topography`. With an augmentation in place that term is $h + r(\mathbf{u} \cdot \hat{\mathbf{n}} - g)$ rather than h , and the difference between the two is measured below.

2.d. Rotating the degrees of freedom

Stop asking for the constraint and impose it. At each constrained node, change the basis in which the velocity unknowns are expressed, from the global Cartesian frame to the local $(\hat{\mathbf{n}}, \hat{\mathbf{t}})$ frame. In that basis “no flow through the boundary” is again a single component, and it is removed the same way it would be on a box.

Collect the per-node rotations into a block-diagonal Q , equal to the identity at every node that is not constrained. The rotated system is

$$\hat{A} = Q^T A Q, \quad \hat{\mathbf{b}} = Q^T \mathbf{b}, \quad \mathbf{u} = Q \hat{\mathbf{u}}, \quad (6)$$

and the wall-normal row of $\hat{A}$ is struck out. The constraint then holds to machine precision, because it is not being solved for at all.

This is the classical answer, not a new one. It is in the early finite-element literature, and Engelman et al. (1982) were already reviewing the alternatives and choosing between them on grounds of global mass conservation in 1982. What is worth explaining is not the idea but why, given that it is exact and the others are not, it is the least used of the three.

3. WHAT IT COSTS, AND THE COST IS STRUCTURAL

Rotating the degrees of freedom leaves the discrete problem in a **mixed basis**. Interior nodes hold (u_x, u_y) ; constrained nodes hold (u_n, u_t) . Nothing about that is difficult in itself, and everything downstream has to agree about which nodes are which.




Figure 1: Where the rotation lives. The obligation is contained: the velocity solve is rotated and carries its multigrid with it, while the Schur complement and the pressure solve beside it never handle a rotated vector, because the pressure block carries no boundary condition of this kind. One un-rotation sits on the boundary between them and feeds both.

Four things carry Q : the operator, the right-hand side, the solution on the way out, and the multigrid prolongation. The coarse operators inherit it through the Galerkin triple product rather than being rotated separately, and the coarse solve then has to be an SVD, because a Galerkin-coarsened rotated operator inherits the rigid-rotation null space of the constrained problem and a redundant LU factorisation meets a zero pivot in it.

That last point is worth dwelling on, because it is the one that surprises. The constraint never has to be re-imposed on a coarse mesh — which is fortunate, since we install no discretisation there and could not impose it if we wanted to. It arrives anyway, algebraically, through $P^T \hat{A} P$ with a rotated P . The evidence that it really arrives is that the coarse operator inherits the null space, which is a property only the constrained problem has.

None of this is an argument against the method. It is an argument for knowing what is being taken on, and it is the honest reason a weakly imposed condition survives in codes that could do this instead.

4. WHICH NORMAL

A question with a less obvious answer than it looks, and the measurements further down say it is the most consequential choice in the note. The boundary of a discretised domain is a set of straight facets, and the assembled constraint is an integral over those facets. The node normal consistent with that integral is the average of the adjacent facet normals **weighted by facet measure** — not the normal of the smooth surface the mesh approximates, and not the facet normal on its own.

This is the consistent normal of Engelman et al. (1982), and we should say so plainly: we re-derived it from the assembled boundary integral, and Engelman, Sani and Gresho derived it in 1982 from global conservation of mass, which is the same object arrived at from the other side. Their paper is about exactly this — how to impose a normal or tangential condition on a boundary that does not line up with the coordinate directions — and the recommendation is the one here.

Using the analytic normal is exact for the geometry and therefore inconsistent with the discretisation, which is the wrong way round: the solver is not solving on the sphere, it is solving on the polyhedron. Using the facet normal is consistent with each facet separately and over-constrains the nodes between them, which is

the failure measured in “A constraint that is satisfied, and wrong” below. In parallel the same argument has a sharper edge, because a normal accumulated rank-locally is wrong at a partition boundary, where a node’s facets are split across ranks and no rank sees them all.

The consistent normal is not the end of the matter. Behr (2004) takes it as the starting point — “preferred from the point of view of conservation” — and reports that in sloshing problems it still does not guarantee a good discrete slip condition, with non-physical recirculation appearing at curved walls; the remedies offered there are the Navier slip condition and a “BC-free” boundary. We have read the abstract rather than the paper, and have not looked for that recirculation in our own cases. It is the obvious thing to test next for anyone running free slip on a strongly curved wall.

5. THE OPTION THIS NOTE LEAVES OUT

Solving in spherical or cylindrical components makes the wall-normal direction a coordinate direction again, and the constraint returns to being “hold one component”. That is the same rotation as above, applied once for the whole domain instead of node by node, and applied that way it costs nothing structurally: there is no mixed basis, because every node is in the same basis.

It works exactly when the boundary lies along a coordinate surface. A sphere, an annulus, a cylinder. It does nothing for topography, for a mesh that has been deformed, or for a tilted internal surface, which is the general case this note is about. The per-node rotation is what remains once the geometry stops cooperating, and its structural price is what buys the generality.

6. WHEN THE CHOICE MATTERS

Here is the awkward part. Solve a convection model with any of these and the velocity field is the same to plotting accuracy, as long as each is set up properly. A leak of order 10^{-3} in $\mathbf{u} \cdot \hat{\mathbf{n}}$ is invisible to anything that consumes the velocity, and consuming the velocity is most of what a model does. If that is your situation, use the simplest thing that works — and read the section on which normal before you do, because that is the choice that can spoil the velocity as well.

The difference appears when the wall-normal traction is the answer rather than a by-product: dynamic topography, a plate-boundary force balance, anything compared against a geoid or a gravity field. That is what the two sections after the leak measure, against an exact surface stress, and the answer is not the one we expected. The treatments that impose the constraint properly all recover the same surface stress at a given mesh, because the recovery sets the floor rather than the boundary condition; what the rotated constraint buys is a route to the traction that does not go through a recovery at all, and it is three times better on the same solve.

Under the rotated constraint the reaction is σ_{nn} : it is the multiplier the solve has already computed, and it comes out of `boundary_normal_traction` without differentiating the answer.

6.a. The leak, measured

An annulus, no slip on the inner radius, the treatment under test on the outer, driven by a degree-four radial density anomaly. The number is the largest normal velocity on the outer boundary, taken against the true radial direction and divided by the flow speed: the fraction of the flow going through a boundary nothing should pass through. Penalty at 10^4 , Nitsche at $\gamma = 10$.

The penalty appears twice, because the normal it is written against turns out to matter more than the method it is written into. One column uses the quadrature-point facet normal `mesh.Gamma`, which is the documented form; the other uses the measure-weighted node normal `mesh.boundary_normal`, which is what Nitsche and the rotated constraint use by default.

cell size	penalty, facet normal	penalty, node normal	Nitsche	multiplier	rotated
0.150	4.5×10^{-3}	3.3×10^{-3}	4.6×10^{-3}	8.3×10^{-4}	7.3×10^{-11}
0.100	2.5×10^{-3}	2.7×10^{-3}	2.2×10^{-3}	1.6×10^{-4}	6.5×10^{-11}
0.075	8.5×10^{-4}	2.6×10^{-3}	1.7×10^{-3}	5.6×10^{-5}	1.0×10^{-10}
0.050	9.5×10^{-4}	2.6×10^{-3}	5.8×10^{-4}	1.2×10^{-5}	8.4×10^{-11}

Nitsche leaks parts in a thousand and improves roughly as $h^{1.9}$ — the rate consistency buys. The multiplier starts an order of magnitude better and falls much faster, near $h^{3.9}$, because the constraint row is an equation rather than a term whose weight has to be chosen. The rotated constraint does not move at all: it sits at the solver’s floor at every resolution, because the mesh has nothing to do with it. The penalty against the node normal does not improve with the mesh either, and for the opposite reason: at a fixed coefficient its leak is set by the coefficient.

The control matters more than the result. With the outer boundary left free the same measurement reads **0.98** — nearly all the boundary flow is normal — so the metric can see a leak when there is one.

The first column is the one to be careful with. Its leak falls with the mesh, which reads as the method working, and it is not: the surface stress measured further down says that solve is 60% wrong in the velocity and 26% wrong in the stress, and refining it does not help. The leak is small because the boundary is being frozen. A metric that only asks whether the constraint is satisfied cannot tell those apart.

6.b. What each parameter buys

The two weak methods look alike in that table. They are not alike, and their own parameters are what tells them apart.

penalty, facet	leak	penalty, node	leak	γ (Nitsche)	leak
10^2	2.6×10^{-1}	10^2	2.6×10^{-1}	1	diverged
10^3	2.3×10^{-2}	10^3	2.6×10^{-2}	10	1.7×10^{-3}
10^4	8.5×10^{-4}	10^4	2.6×10^{-3}	100	2.7×10^{-4}
10^5	9.6×10^{-4}	10^5	3.0×10^{-4}	1000	3.0×10^{-5}
10^6	diverged	10^6	4.5×10^{-5}	10^4 and above	diverged

Nitsche is bounded at both ends. Below $\gamma = 1$ the form is no longer coercive and no amount of tuning recovers it; from $\gamma = 10^4$ in this problem the line search stops converging. The virtue of $\gamma = 10$ is that it sits in the middle of that window on any mesh, because γ is dimensionless and the term it scales already carries μ/h . The penalty coefficient carries no such scaling, which is why the value that works is a property of the problem rather than a default — and, on the second half of this note’s test, of the local viscosity as well.

The two penalty columns are the same method against different normals, and they behave differently in a way the leak alone does not explain. Against the node normal the leak keeps falling, a decade of coefficient for a

decade of leak, all the way to 10^6 . Against the facet normal it stops improving after 10^4 and the solve fails at 10^6 . That looks like the conditioning wall a penalty is expected to have. It is not: what stalls is the leak, because by 10^4 the boundary is nearly frozen and there is little normal flow left to remove. The stress section below measures the same runs against an exact answer, and they are getting worse throughout.

6.c. The stress, measured

The leak says how well each treatment holds the boundary. Whether the answer is right is a different question, and it needs an exact answer to compare against.

Kramer, Davies and Wilson (Kramer et al., 2021) give exact Stokes solutions in a cylindrical annulus, and their `assess` package publishes the radial stress itself rather than leaving it to be recovered from a velocity field. Underworld wraps it as `uw.analytic.CylindricalStokes`. The case here is the smooth one: a density anomaly $(r/r_o)^k \cos n\theta$ with $n = 2$ and $k = 3$, viscosity 1, free slip on both radii. On the outer boundary the exact radial stress is a single harmonic,

$$\sigma_{rr}(r_o, \theta) = 0.1506696 \cos 2\theta, \quad (7)$$

fitted to a residual of 10^{-16} , so the whole of the surface stress is that one amplitude and the metric is its relative error.

The inner boundary carries the exact velocity as a Dirichlet condition instead of a free-slip treatment of its own. The exact solution satisfies both, so the problem is unchanged, and the treatment under test is then the only free-slip condition in the model.

There are two routes to the surface stress and the difference between them is the point of the section:

- **recovered** — project $\hat{n} \cdot \sigma \hat{n}$ out of the solved velocity and pressure. Every treatment can do this, and it is the only route the weak ones have.
- **reaction** — `boundary_normal_traction` for the rotated constraint, and the multiplier field for the constraint method. Not recovered from the solution: an unknown the solve returned.

Both are taken against the true radial direction, which is also the direction the oracle publishes, so no treatment is scored against its own normal. The reaction and the multiplier come back with the opposite sign to σ_{rr} — they are the traction holding the boundary rather than the traction the fluid exerts, and `dynamic_topography` carries the sign back — so amplitudes are compared unsigned.

cell size	penalty, facet	penalty, node	Nitsche	constraint	rotated	rotated, reaction	multiplier field
0.150	1.5×10^{-1}	2.5×10^{-2}	5.9×10^{-2}	2.4×10^{-2}	2.4×10^{-2}	6.8×10^{-3}	8.6×10^{-3}
0.100	1.3×10^{-1}	1.1×10^{-2}	2.4×10^{-2}	1.0×10^{-2}	1.0×10^{-2}	3.3×10^{-3}	8.5×10^{-4}
0.075	1.1×10^{-1}	7.2×10^{-3}	1.5×10^{-2}	6.3×10^{-3}	6.2×10^{-3}	2.1×10^{-3}	1.7×10^{-3}
0.050	6.3×10^{-2}	3.6×10^{-3}	6.3×10^{-3}	2.7×10^{-3}	2.7×10^{-3}	1.1×10^{-3}	1.4×10^{-3}

The first five columns are all the recovered stress, so they differ only by which boundary condition produced the field. The last two are the reaction and the multiplier on the same solves. Penalties at 10^4 , Nitsche at $\gamma = 10$, which are the values the leak table used. With the outer boundary left free the same measurement reads 1.0 — the surface stress is gone — so the metric can see the condition being removed.

Once the constraint is imposed against the node normal, which treatment imposed it stops mattering to the recovered stress. At cell 0.075 the rotated constraint gives 6.2×10^{-3} , the multiplier 6.3×10^{-3} , the penalty at 10^5 6.3×10^{-3} and Nitsche at $\gamma = 100$ 6.6×10^{-3} . Those are the same number. What sets it is the recovery — a projection of a stress differentiated out of a piecewise-quadratic velocity — and not the boundary condition underneath. The reasoning this note began with — that a traction recovered from an approximate constraint inherits the approximation — is not what the measurement shows once the constraint is written against the right normal. It shows a floor that all of them share.

The reaction is about three times better than the recovery, on the same solve. 2.1×10^{-3} against 6.2×10^{-3} at cell 0.075, and it converges at the same rate rather than a better one. It costs nothing — it is the residual the solve has already assembled, and no field is differentiated to get it — and the timings below put a number on “nothing”.

The multiplier’s traction is the same story from the other side, and lands in the same place: 8.5×10^{-4} to 1.7×10^{-3} over the three finer meshes. Its column does not fall smoothly with h , because part of what it reports is the augmentation times the constraint residual, and how far a particular solve drove that residual is not a function of the mesh.

The parameter that was enough for the leak is not enough for the stress. Nitsche at $\gamma = 10$ leaks 1.2×10^{-3} in this problem and gets the stress amplitude 1.5% wrong. At $\gamma = 100$ the leak improves by a factor of nearly forty and the stress by a factor of two, to the recovery floor, where more γ buys nothing. Reading the leak alone would have said $\gamma = 10$ was converged.

6.d. *A constraint that is satisfied, and wrong*

The facet-normal penalty column does not converge. That is not a slow rate; the error grows slightly as the mesh is refined, and the leak is excellent throughout.

cell size	leak	velocity error	stress error
0.150	8.9×10^{-6}	0.60	0.21
0.100	8.2×10^{-6}	0.61	0.24
0.075	1.1×10^{-5}	0.61	0.25
0.050	1.9×10^{-5}	0.60	0.26
0.035	2.3×10^{-5}	0.59	0.26




Figure 2: The same problem, the same coefficient of 10^6 , the same colour scale, and the only difference between the middle and right panels is which normal the penalty is written against. Against the facet normal the flow along the outer boundary is suppressed: the peak speed falls by a quarter and the two fast lobes at the boundary are gone. Refining the mesh does not bring them back.

The coefficient is 10^6 and the normal is `mesh.Gamma`, the facet normal at the quadrature points. The same coefficient against the measure-weighted node normal, on the same meshes, gives velocity errors of 1.0×10^{-2} , 2.4×10^{-3} , 1.0×10^{-3} and 4.9×10^{-4} and stress errors of 2.4×10^{-2} , 6.3×10^{-3} , 2.7×10^{-3} and 1.4×10^{-3} . One difference, one line of code, and one method converges while the other does not.

The mechanism is old and well understood. Imposing $\mathbf{u} \cdot \hat{\mathbf{n}} = 0$ facet by facet asks a node shared by two facets to satisfy two different constraints, and two independent constraints on a two-component velocity leave nothing. Push the coefficient up and the vertex velocities go to zero: the flow is being asked to stay inside a polygon rather than a circle, and the polygon is not a good enough approximation of the circle for that particular question. The discrete limit is a different problem from the smooth one, so refining the mesh does not approach the smooth answer.

Two things follow for practice. A direct penalty is a perfectly good method, but it has to be written against the node normal. And the leak is not a sufficient check: at a coefficient of 10^3 the facet-normal penalty leaks 3×10^{-2} and gets the stress amplitude right to 1.9×10^{-3} , while at 10^8 it leaks 1×10^{-7} and is 26% wrong. Over that range the constraint improves by five orders of magnitude and the answer gets steadily worse.

6.e. The multiplier was not the whole traction

The multiplier column of the table above used to sit near 3×10^{-2} for three resolutions and then drop, which is not a convergence rate. The cause was the augmented-Lagrangian term, and it was a defect rather than a property of the method.

The momentum row carries both the multiplier and the augmentation, $(h + r(\mathbf{u} \cdot \hat{\mathbf{n}} - g))\hat{\mathbf{n}}$, so the traction holding the boundary is $h + r(\mathbf{u} \cdot \hat{\mathbf{n}} - g)$ and not h alone. The second term vanishes only where the constraint row is satisfied exactly. Discretely it is satisfied to the solver's tolerance, and r multiplies that residual straight back into the traction. Underworld returned h .

On the annulus, with a uniform viscosity and the default $r = 10^4\mu$, the omitted share is a few per cent. On SolCx it is nearly everything, because the default r is viscosity-weighted and the step takes it to 10^{10} on the stiff half. At a contrast of 10^6 , against an exact surface topography whose deviation peaks at 0.381:

read	peak	correlation with the exact	relative l_2
the multiplier h	0.042	-0.52	0.96
$h + r(\mathbf{u} \cdot \hat{\mathbf{n}} - g)$	0.382	+0.999	0.084

An order of magnitude low and pointing the wrong way, on a solve whose velocity error is 8.8×10^{-6} . Underworld3 issue #607: `traction()` now returns the sum and `topography()` is built on it, so the tables here are the corrected read and a caller who asks for topography gets it.

Worth saying how it survived. The method's own validation scored a *correlation* of 0.9999 between the multiplier and the recovered normal stress. A correlation is scale-free, so it cannot see a systematic amplitude deficit – which is exactly what a missing share of the load is. The replacement guard scores a relative l_2 against an exact answer and carries the bare multiplier as its negative control.

6.f. The multiplier and the consistent boundary flux are the same object

The correction above is not a patch. Write the momentum row's boundary term out and the identity is immediate: the assembled load is $M_\Gamma (h + r(\mathbf{u} \cdot \hat{\mathbf{n}} - g))$, and at convergence it balances the volume residual restricted to the boundary, which is precisely the nodal load the consistent boundary flux back-calculation reads (Zhong et al., 1993). So

$$h + r(\mathbf{u} \cdot \hat{\mathbf{n}} - g) = -M_\Gamma^{-1}(A\mathbf{u} - \mathbf{b})|_\Gamma, \quad (8)$$

the CBF traction de-smearred with the boundary mass. The multiplier is not a second, independent estimate of the surface stress: it is the same computation the rotated constraint's reaction performs, arrived at by carrying the traction as an unknown instead of reading it out of the residual afterwards. Dropping the r term is dropping part of the load, which is why it fails exactly where r is large.

Measured across the two solves — they are different discrete problems, so this is agreement rather than an identity check — the corrected multiplier and the rotated reaction differ by 3.2% at a contrast of 100 and 4.9% at 10^6 , both inside each route's own error against the exact answer (5 to 9%). They cannot be compared within a single solve, and the reason is the identity itself: on a multiplier-constrained boundary the constraint enters the row it constrains, so the assembled residual there is balanced at convergence and the CBF back-calculation reads zero (measured: rms 4×10^{-13} against a traction of 0.37). There is no reaction left to read because the multiplier is holding it. `boundary_flux` now says so instead of returning that zero quietly (underworld3#614).

This also settles a question the free-surface work left open. That work used SolCx the same way, to choose among topography recoveries, and landed on a rotated free-slip lid with the CBF reaction and a continuous pressure — corr 0.999, relative l_2 0.04 — while rejecting the multiplier. Both conclusions were right about what was in front of them: the multiplier *as returned* is missing the augmentation share, and the CBF reaction is the same quantity with nothing missing.

6.g. The other half: a lateral viscosity contrast

No exact solution has both a curved boundary and a laterally varying viscosity, so the case where weak constraints are most often reported to give trouble is a separate test with a trivial geometry. SolCx is that test: the unit box, free slip on all four walls, viscosity 1 to the left of $x = 0.5$ and η_B to the right. `uw.analytic.SolCx` publishes the exact dynamic topography on the top wall. Three walls carry the ordinary component condition and the treatment under test is on the top wall alone.

On a box every treatment reduces to holding one velocity component, so nothing here is about normals. What it can say is whether a treatment holds the traction it was given when the viscosity beside it jumps.

Relative l_2 error of the surface topography along the top wall, mean removed, at 32×32 elements. Each entry is the whole wall and then the wall with two elements trimmed from each end.

η_B/η_A	component Dirichlet	penalty, 10^4	multiplier	rotated
10	0.048 / 0.054	0.045 / 0.051	0.048 / 0.054	0.048 / 0.054
10^2	0.072 / 0.081	0.056 / 0.060	0.072 / 0.081	0.072 / 0.081
10^3	0.075 / 0.084	0.234 / 0.230	0.075 / 0.084	0.075 / 0.084
10^4	0.076 / 0.085	0.698 / 0.697	0.076 / 0.085	0.076 / 0.085
10^6	0.076 / 0.085	0.992 / 1.000	0.075 / 0.084	0.076 / 0.085

The three exact treatments agree to three figures at every contrast, whole wall and trimmed alike. That is the result to take from this half, and it took two fixes to get: the multiplier had to report the whole traction rather than h , and the rotated constraint had to hold the corner where it meets the side walls. Before them the rotated column read 0.322 at a contrast of 10^{-6} — six times the reference, all of it two nodes — and the multiplier column was an order of magnitude out at 10^6 .

Read the first column as the floor. The component Dirichlet condition is exact and has no parameter, and its velocity error is 8.8×10^{-6} at a contrast of 10^6 . It still reads 0.085. That number is the recovery's error, not a boundary condition's: on the stiff half the recovered σ_{zz} is a difference between the pressure and $2\eta \partial_z u_z$ with $\eta = 10^6$, so a relative velocity error of 10^{-5} arrives in the stress at the size of the signal. Nothing here is worse than the reference except the penalty.

A bare penalty coefficient cannot serve both halves. At 10^4 it is the best column in the table at low contrast — the constraint is weak enough not to fight the recovery — and by 10^6 it is useless: 0.992, which is to say the recovered topography carries none of the signal. Scaling the coefficient by the local viscosity is the obviously right thing to want, and it does not solve here at any magnitude we tried, from μ to $10^3 \mu$ (a `Piecewise` viscosity inside the boundary term fails the line search). This is the same lesson as γ 's window on the annulus, in a place where the window closes entirely.

Nitsche is missing from this table, and honestly so. Our configuration of it on this box converges at a contrast of 10^6 and fails the line search at 10^{-3} — the opposite way round from every expectation — at 16×16 and 32×32 alike and at $\gamma = 10, 100$ and 1000 . Where it does converge, γ has to rise with the contrast exactly as the annulus said: at 10^6 and 16×16 the surface stress error is 25 at $\gamma = 10$, 1.2 at 100 and 0.17 at 1000, while the constraint is held to 10^{-3} or better throughout. We are not confident enough in that configuration to put a column of numbers behind it.

A measurement we withdrew

An earlier draft of this section carried a Nitsche column and different numbers throughout. They were taken while several runs shared one mesh-cache file: `StructuredQuadBox` keys its cache on the box corners and not on the element resolution, so concurrent runs at different resolutions silently swap meshes (underworld3#618). A marginal solve then flips between converged and diverged for reasons that look like the method. Every number in this section was re-measured sequentially, on a cleared cache, with nothing else running.

6.h. What each one costs

Accuracy is half the choice. Seconds on the annulus, median of three timed repeats after an untimed warm-up, run sequentially — the solve, and then the recovery of the surface traction by whatever route that treatment has.

cell size	velocity nodes	penalty	Nitsche	multiplier	rotated
0.075	2 174	0.01 / 0.153	0.01 / 0.037	0.02 / 0.001	0.01 / 0.003
0.050	4 802	0.03 / 0.058	0.03 / 0.060	0.04 / 0.002	0.03 / 0.004
0.035	9 406	0.06 / 0.095	0.06 / 0.099	0.08 / 0.004	0.05 / 0.006

The solve costs the same whichever you choose. Within the spread, all four are the same number at every size — 0.06 s at 9 400 nodes. The multiplier is the one that carries a visible premium, about 30%, which is the

extra field and the larger saddle point; its interior degrees of freedom are constrained out of the global system, so what is left is the boundary trace.

The recovery is where they differ, by a factor of 15 to 25. Projecting $\hat{n} \cdot \sigma \hat{n}$ is a second solve, and it costs more than the Stokes solve did at the coarser sizes. The reaction and the multiplier are arithmetic on the boundary trace: 4 to 6 ms where the projection takes 95 to 99.

That ratio is the practical argument, and it points the same way the accuracy did. The route that does not differentiate the answer is both the more accurate one and the cheaper one, and on a time-dependent model it is paid at every step.

What these timings are not

Two-dimensional problems of ten thousand nodes, on one core, with a direct solver. They separate a recovery solve from boundary arithmetic, which is a structural difference and will survive scaling. They say nothing about how the four compare on a large parallel spherical shell, where the rotated velocity block's multigrid and the multiplier's larger Schur complement are the terms that matter and neither is exercised here.

6.i. Which one to use

For a model that consumes the velocity and nothing else, all four are the same to plotting accuracy — provided the constraint is written against the node normal. That proviso is the only one that can spoil the velocity, and it costs one line.

When the wall-normal traction is the answer:

- **Rotated free slip is the default.** It holds the constraint to machine precision rather than to the discretisation, its reaction is the most accurate surface stress measured here, and that reaction is nearly free. The price is structural — a mixed basis that the multigrid has to carry — and it is paid once, inside the solver, rather than by the person setting up the model.
- **The multiplier is its equal on accuracy** and returns the same object by a different route; take it when you want the traction as an unknown of the system, or when the constrained problem's conditioning suits you better. It costs about 30% more to solve.
- **Nitsche is the one to reach for when the boundary condition must change during the model** — a wall that begins as a prescribed velocity and relaxes to a prescribed traction is a Nitsche problem, because a hard constraint cannot morph. Budget for tuning γ against the stress and not against the leak, and expect the window to move with the viscosity contrast.
- **A direct penalty is fine for a velocity-only model** and needs the node normal, a coefficient chosen per problem, and a check on something physical before the answer is believed.

7. USING IT

```
import underworld3 as uw

mesh = uw.meshing.Annulus(radiusInner=0.5, radiusOuter=1.0, cellSize=0.05)
stokes = uw.systems.Stokes(mesh)

# Value first: 0 is free slip. A non-zero scalar or expression prescribes the
```

```
# wall-normal datum u.n = u_n strongly instead.
stokes.add_rotated_freeslip_bc(0.0, "Upper")
stokes.add_rotated_freeslip_bc(0.0, "Lower")

stokes.solve()

# The constraint reaction, which is the boundary normal traction.
sigma_nn = stokes.boundary_normal_traction("Upper")
```

Leave the normal to Underworld unless the constraint has to follow the true surface rather than the mesh. Passing an analytic normal — $\mathbf{x} / |\mathbf{x}|$ on a sphere — is exact for the geometry and keeps a consistency error against the faceted assembly, which is usually not what you want.

Reach for Nitsche when the boundary condition has to **change during the model**. A hard constraint cannot morph: a wall that begins as a prescribed velocity and relaxes to a prescribed traction is a Nitsche problem, because the rotated constraint is either imposed or it is not.

REFERENCES

- Behr, M. (2004). On the application of slip boundary condition on curved boundaries. *International Journal for Numerical Methods in Fluids*, 45(1), 43–51. <https://doi.org/10.1002/fld.663>
- Engelman, M. S., Sani, R. L., & Gresho, P. M. (1982). The implementation of normal and/or tangential boundary conditions in finite element codes for incompressible fluid flow. *International Journal for Numerical Methods in Fluids*, 2(3), 225–238. <https://doi.org/10.1002/fld.1650020302>
- Kramer, S. C., Davies, D. R., & Wilson, C. R. (2021). Analytical solutions for mantle flow in cylindrical and spherical shells. *Geoscientific Model Development*, 14(4), 1899–1919. <https://doi.org/10.5194/gmd-14-1899-2021>
- Nitsche, J. (1971). "Über ein Variationsprinzip zur Lösung von Dirichlet-Problemen bei Verwendung von Teilräumen, die keinen Randbedingungen unterworfen sind. *Abhandlungen Aus Dem Mathematischen Seminar Der Universität Hamburg*, 36(1), 9–15. <https://doi.org/10.1007/bf02995904>
- Zhong, S., Gurnis, M., & Hulbert, G. (1993). Accurate determination of surface normal stress in viscous flow from a consistent boundary flux method. *Physics of the Earth and Planetary Interiors*, 78(1–2), 1–8. [https://doi.org/10.1016/0031-9201\(93\)90078-n](https://doi.org/10.1016/0031-9201(93)90078-n)